

# Mastic: Private Weighted Heavy-Hitters and Attribute-Based Metrics

Tech Talk @ Google – 2024/10/3

Dimitris Mouris (Nillion)

**Christopher Patton (Cloudflare)**

Hannah Davis (Seagate)

Pratik Sarkar (Supra Research)

Nektarios G. Tsoutsos (University of Delaware)

# Work at IETF

- Privacy Preserving Measurement (PPM)
  - Use cases
    - aggregation (à la Prio): average, variance, histograms, Bloom filters, linear regression, step of gradient descent
    - heavy hitters (à la Poplar): compute the set of the most frequent inputs (e.g., the most visited websites)
  - Privacy: distribute computation among multiple servers so that no one server sees any measurement in plaintext
  - Robustness: detect invalid measurements uploaded by misbehaving clients
- Verifiable Distributed Aggregation Functions (VDAFs): Specifications of Prio and Poplar
- Distributed Aggregation Protocol (DAP): Execution of a VDAF over HTTP

# Where does Mastic fit?

- There are things we'd like Poplar or Prio to do:
  - *attribute-based* aggregation: average, variance, histograms, Bloom filters, linear regression, step of gradient descent *grouped by client attributes (e.g., user agent or geolocation)*
  - *weighted* heavy-hitters: compute the set of the ~~most frequent~~ *heaviest weight* inputs (e.g., the ~~most visited websites~~ *websites with the highest engagement*)
- Poplar has rough edges
  - Two rounds of communication for input validation (MPC multiplication w/ Beaver triples provided by clients)
  - Different field for leaf versus inner nodes of the IDPF tree, unsafe to use intermediate outputs

# Design goals

- Ease of implementation: simplify design, reuse existing components where possible
- Deployability: one round of communication between aggregation servers instead of two
- Squeeze more functionality out of less code: *attribute-based* aggregation and *weighted* heavy-hitters

# Agenda

- 1 Setting (functionality, architecture, security goals) ~10 minutes
- 2 Primitive #1: VIDPF ~15 minutes
- 3 Primitive #2: FLP ~3 minutes
- 4 Putting it all together ~3 minutes
- 5 Next steps ~5 minutes

# Setting

# Functionality

- Each client generates a **measurement** consisting of an **input** (an element of  $\{0,1\}^n$ ) and its **weight** (an element of  $\mathbb{F}^m$ )
- For each **prefix**, compute the **total weight** of all inputs beginning with the prefix

```
def mastic_func(measurements: list[tuple[Index, Weight]],
               prefixes: list[Index]) -> dict[Index, Weight]:
    """
    Compute the total weight for each prefix for the set of
    measurements.
    """
    r: dict[Index, Weight] = {}
    for (alpha, beta) in measurements:
        for x in prefixes:
            if x.is_prefix(alpha):
                total_weight = r.setdefault(x, Weight(0))
                total_weight += beta
    return r
```

# Functionality

```
def weighted_heavy_hitters(measurements: list[tuple[Index, Weight]],
                           threshold: Weight,
                           bit_len: int) -> list[Index]:
    """
    Compute the weighted heavy hitters for the given threshold.
    """
    prefixes = [Index(0), Index(1)]
    for level in range(bit_len):
        next_prefixes = []
        for (x, total_weight) in mastic_func(measurements, prefixes).items():
            if total_weight >= threshold:
                if level < bit_len - 1:
                    next_prefixes.append(x.left_child())
                    next_prefixes.append(x.right_child())
                else:
                    next_prefixes.append(x)
        prefixes = next_prefixes
    return sorted(prefixes)
```

Weighted heavy hitters  
uses Mastic as a  
subroutine over multiple  
rounds of aggregation



## Use case #1: Ad attribution

- ***On-device attribution***: browser keeps track of which ads are shown to the user: when the user makes a purchase, the browser determines which ad, if any, deserves credit.
  - Input: a unique identifier for the ad campaign (e.g., "Bartlet for America")
  - Weight: whether a purchase was made (Count); how much the user spent (Sum); which category of product was purchased (Histogram); etc.

## Use case #2: Browser telemetry

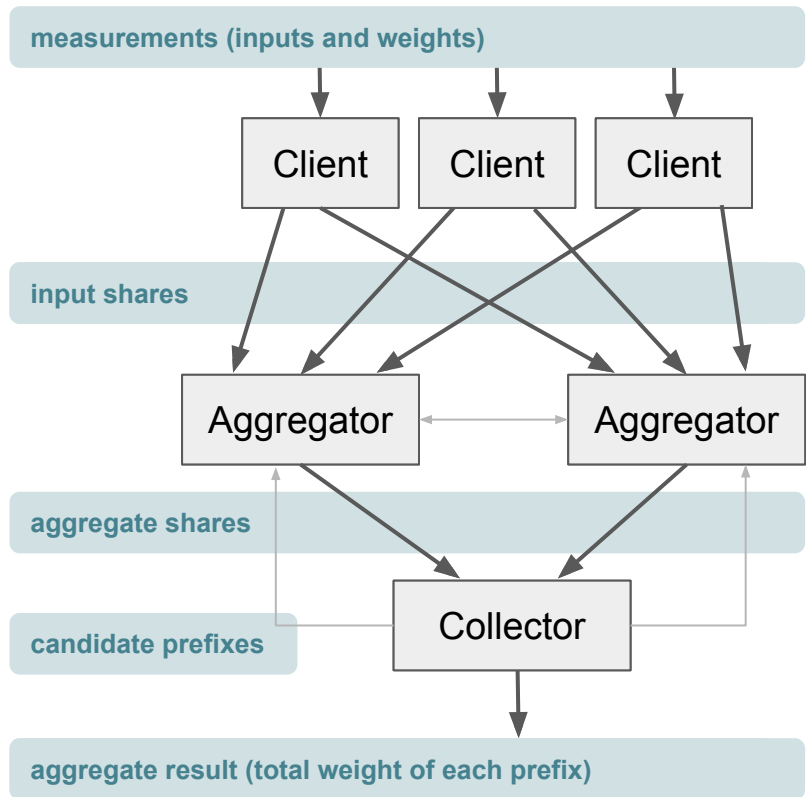
- Many use cases for Prio-style metrics in browsers, e.g.: page load metrics for a fixed set of "benchmark" websites (e.g., [google.com](https://www.google.com) or [zombo.com](https://www.zombo.com))
  - Input: encodes whatever attributes we might want to break down by, eg., software version and geographic location
  - Weight: average ([Sum](#)) or distribution ([Histogram](#)) of load time

## Use case #3: Network error logging

- NEL: To detect configuration issues, content-delivery networks like Cloudflare want to know the types of errors preventing eyeballs from connecting to their customers (DNS, TCP, TLS, etc.)
  - Input: the customer identified by a DNS name
  - Weight: indication of which error occurred (**Histogram**)

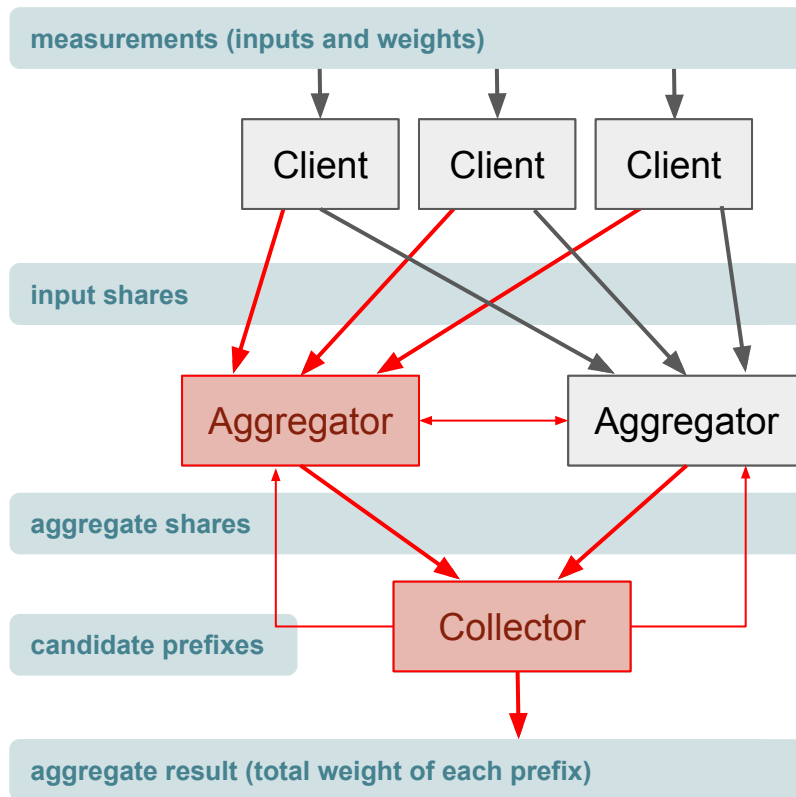
# Architecture

- Clients **shard** their measurements
- Aggregators interact to validate the measurements, then compute shares of the total weight for each candidate prefix
- Collector **unshards** the aggregate result



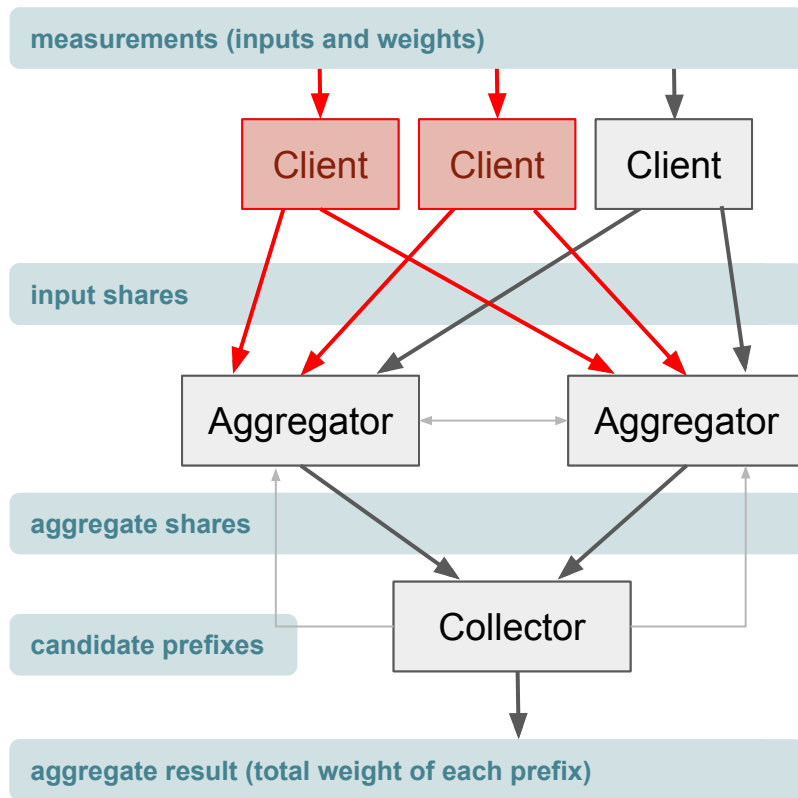
# Security goal #1: Privacy

- Attacker controls one of the Aggregators and the Collector
- Attacker's view is efficiently *simulatable* given only aggregate results



## Security goal #2: Robustness


- Attacker controls a subset of the clients and eavesdrops on the network
- Aggregate result computed by the collector is efficiently **extractable** from the input shares



# Primitive #1: VIDPF

# Verifiable Incremental Distributed Point Function (VIDPF)

- **point function:**  $F(x) = \beta$  if  $x = \alpha$  else  $0$
- **incremental:**  $F(x) = \beta$  if  $x.is\_prefix(\alpha)$  else  $0$
- **distributed:** client generates secret shares of  $F$  such that each aggregator computes a secret share of  $F(x)$  for any  $x$  without learning  $\alpha$  or  $\beta$
- **verifiable:** aggregators verify that, for any sequence of prefixes  $prefixes$ ,  $F(x) \neq 0$  for at most one  $x$  in  $prefixes$  and the others are  $0$  (**onehot**)






# Verifiable Incremental Distributed Point Function (VIDPF)

Client:  $(\text{correction\_words}, [\text{key0}, \text{key1}]) = \text{vidpf.gen}(\alpha, \beta)$


Aggregator 0:  $(\text{out0}, \text{proof0}) = \text{vidpf.eval}(0, \text{correction\_words}, \text{key0}, \text{prefixes})$

Aggregator 1:  $(\text{out1}, \text{proof1}) = \text{vidpf.eval}(1, \text{correction\_words}, \text{key1}, \text{prefixes})$

**verifiability:** It's computationally infeasible to find  $\text{prefixes}$  for which  $\text{proof0} == \text{proof1}$  but  $\text{out0} + \text{out1}$  is not onehot.



# A bit of history




[[BG16](#)] Boyle et al. "Function Secret Sharing: Improvements and Extensions." CCS 2016.

[[BBCG+21](#)] Boneh et al. "Lightweight Techniques for Private Heavy Hitters." IEEE S&P 2021.

[[CP22](#)] de Castro and Polychroniadou. "Lightweight, Maliciously Secure Verifiable Function Secret Sharing." Eurocrypt 2022.


[[MST24](#)] Mouris et al. "PLASMA: Private, Lightweight Aggregated Statistics against Malicious Adversaries." PETS 2024.

# The DPF of [BG16]



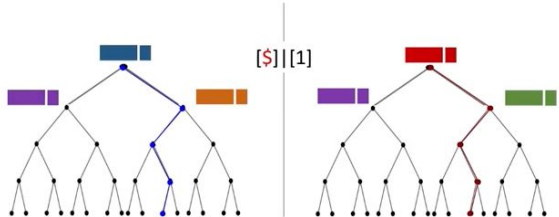
**BIU**  
Center for Research in Applied  
Cryptography and Cyber Security  
Bar-Ilan University

The 12<sup>th</sup> BIU Winter School  
on Cryptography  
Advances in Secure Computation



Prime Minister's Office  
National Cyber Directorate  
National Cyber Bureau

## DPF Construction from PRGs



Invariant for Eval:

For each node  $v$  on evaluation path we have  $[S][b]$

- $v$  on special path:  $S$  is pseudorandom,  $b=1$
- $v$  off special path:  $S=0$ ,  $b=0$

42

Gold Silver

Protocol Labs T11

Sponsored by

STARKWARE


Bronze

Duality

For a complete tutorial, check out [Elette Boyle's series](#) on function secret sharing for the 12th BIU Winter School


# The DPF of [BG16]

- Evaluation of input  $x$ :
  - Initialization:
    - $ctrl$  = aggregator's ID (0 or 1)
    - $seed$  = aggregator's key share
  - for each level  $i$ :
    - Extend the current  $seed$  into a  $ctrl/seed$  for each child node
    - If  $x[i] = 0$  then select the left child; otherwise select the right child.
  - Share of the output is derived from the  $seed$  for the last level



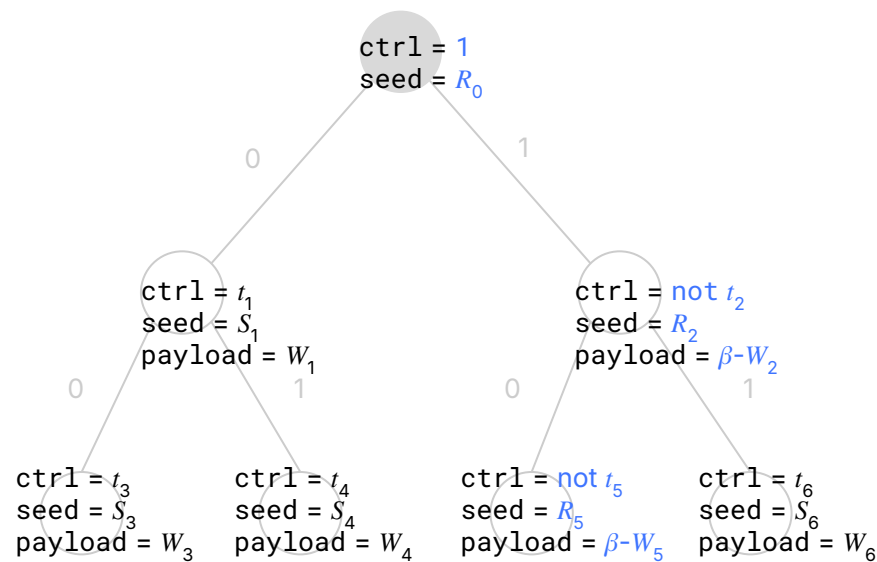
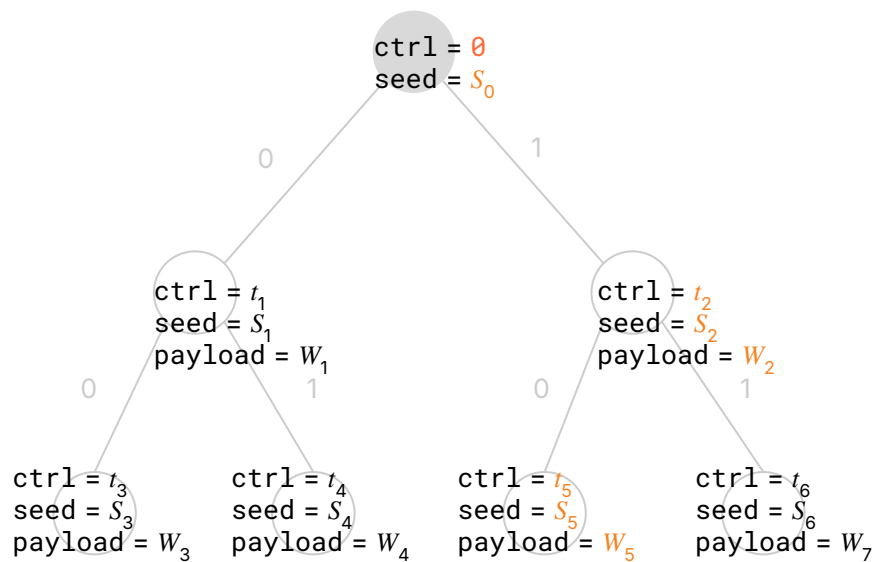
# The DPF of [BGI16]

**DPF invariant:** if  $x[:i].is\_prefix(\alpha)$ , then the **ctrls** are shares of 1 and the **seeds** are distinct; otherwise, the **ctrls** are shares of 0 and the **seeds** are the same. (If **ctrl**==1, then "correct" with the level's **correction word**.)




# DPF → IDPF [BBCG+21]

Generate payloads for intermediate levels: use **seed** to derive the next seed and a share of the payload (after correction).



# DPF → VDPF [CP22]


Verify the DPF invariant: let  $\text{proof} = \text{hash}(\text{seed}, x[:i])$  be the **node proof**; if the DPF invariant holds, then each aggregator computes the same node proof (after correction).



[CP22] de Castro and Polychroniadou. "Lightweight, Maliciously Secure Verifiable Function Secret Sharing." Eurocrypt 2022.

# IDPF, VDPF → VIDPF [MST24]

Generate intermediate payloads, verify the DPF invariant, and verify **path consistency**: the payload of each node traversed is equal to the sum of its children  $\Rightarrow \beta$  is the output for each prefix of  $\alpha$ .




[MST24] Mouris et al. "PLASMA: Private, Lightweight Aggregated Statistics against Malicious Adversaries." PETS 2024.



# IDPF, VDPF → VIDPF [MST24]

Generate intermediate payloads, verify the DPF invariant, and verify **path consistency**: the payload of each node traversed is equal to the sum of its children  $\Rightarrow \beta$  is the output for each prefix of  $\alpha$ .



# Primitive #2: FLP

# Fully Linear Proof (FLP)

Zero-knowledge proof system on distributed data [BBCG+19]. Syntax here is as presented in [draft-irtf-cfrg-vdaf]:

Client:            `proof = flp.prove( $\beta$ );` secret share `proof` into  
                    `proof0 + proof1` and  `$\beta$`  into `beta0 + beta1`

Aggregator 0:   `verifier0 = flp.query(beta0, proof0)`

Aggregator 1:   `verifier1 = flp.query(beta1, proof1)`

***proof verification:*** if `flp.decide(verifier0 + verifier1)`  
then the weight is valid with high probability.

[[BBCG+19](#)] Boneh et al. "Zero-Knowledge Proofs on Secret-Shared Data via Fully Linear PCPs". Crypto 2019.

[[draft-irtf-cfrg-vdaf](#)] Barnes et al. "Verifiable Distributed Aggregation Functions." CFRG draft, version 11.

# Putting it all together

# VIDPF + FLP =



Weight type is determined by the FLP:

- **Count**: weight is 0 or 1 (same functionality as Poplar)
- **Sum**: weight is an integer in a bounded range
- **Histogram**: weight is a histogram with a fixed number of buckets

Client runs the sharding algorithm to get the **public share** (the correction words) sent to both aggregators and one **input share** for each aggregator.

```
def shard(self,
            measurement: tuple[Index, Weight],
            nonce: bytes):
    # Encode the weight according to the FLP.
    ( $\alpha$ , weight) = measurement
     $\beta$  = self.flp.encode(weight)

    # Generate VIDPF keys (depends on the nonce).
    (correction_words, keys) = self.vidpf.gen( $\alpha$ ,  $\beta$ , nonce)

    # Generate FLP and split it into shares.
    proof = self.flp.prove( $\beta$ )
    helper_seed = gen_rand(32)
    helper_proof_share = self.helper_proof_share(helper_seed)
    leader_proof_share = vec_sub(proof, helper_proof_share)

    return (correction_words,
            [(keys[0], leader_proof_share), # aggregator 0
             (keys[1], helper_seed)])      # aggregator 1
```


$$\text{VIDPF} + \text{FLP} = \text{🌿}$$

Aggregator, on input of the report from the client (correcton words, VIDPF key share, FLP proof share) and **prefixes** from the collector:

- Evaluate the VIDPF at **prefixes** to get output **output\_share**, the secret shares of  $F(x)$  for each  $x$  in **prefixes**
- Interact with the co-Aggregator to check (*one roundtrip over the network*)
  - **onehotness**: Check that DPF invariant holds for **prefixes**  $\Rightarrow$  at most one of the outputs is  $\beta$  (the rest are  $0$ )
  - **path consistency**: For each node traversed in the prefix tree, check that the payload is equal to the sum of its children  $\Rightarrow F(x) = \beta$  for each prefix  $x$  of  $\alpha$
  - **weight validity**: Verify the FLP  $\Rightarrow \beta$  encodes a valid weight
- Aggregate **output\_share**.

# Next steps

# Planned work

Status	Task
Done 	Initial design and analysis [MPD+25].
In progress	Implementation ( <a href="https://github.com/divviup/libprio-rs">github.com/divviup/libprio-rs</a> ) in collaboration with ISRG ( <a href="https://divviup.org/">Divvi Up</a> ), who maintains open source implementations of Prio and Poplar (Rust).
In progress	Specification [draft-mouris-cfrg-mastic]. Many improvements are being considered ( <a href="https://github.com/jimouris/draft-mouris-cfrg-mastic/issues">github.com/jimouris/draft-mouris-cfrg-mastic/issues</a> ).
Planned	End-to-end analysis of the final specification, especially concrete security of VIDPF.

[[MPD+25](#)] Mouris et al. "Mastic: Private Weighted Heavy-Hitters and Attribute-Based Metrics." PETS 2025.

[[draft-mouris-cfrg-mastic](#)] Mouris et al. "The Mastic VDAF." Individual draft, version 03.



# Status of Mastic at IETF

- [draft-mouris-cfrg-mastic] is an **individual draft** and has not been adopted by a working group
- We will pursue adoption if/when there is sufficient interest to deploy Mastic. Currently in a holding pattern:
  - Top priority right now is completing the core documents [draft-ietf-ppm-dap, draft-irtf-cfrg-vdaf]
  - CFRG currently has low bandwidth for new work
  - PPM is chartered to do protocol design, not develop new cryptographic algorithms (this is meant to be delegated to CFRG)
  - We considered replacing Poplar with Mastic, but the current consensus is to keep Poplar in the core VDAF draft.

[[draft-mouris-cfrg-mastic](#)] Mouris et al. "The Mastic VDAF." Individual draft, version 03.

[[draft-ietf-ppm-dap](#)] Geoghegan et al. "Distributed Aggregation Protocol for Privacy Preserving Measurement." PPM working group draft, version 11.

[[draft-irtf-cfrg-vdaf](#)] Barnes et al. "Verifiable Distributed Aggregation Functions." CFRG draft, version 11.


# Thank you

 [cpatton@cloudflare.com](mailto:cpatton@cloudflare.com)

 [datatracker.ietf.org/doc/draft-mouris-cfrg-mastic](https://datatracker.ietf.org/doc/draft-mouris-cfrg-mastic)


# Functionality

**prefix tree:** Label each node with the total weight of the prefix given by the path from the root




# Functionality

**prefix tree:** Label each node with the total weight of the prefix given by the path from the root



# Functionality

**prefix tree:** Label each node with the total weight of the prefix given by the path from the root



**prefix tree:** Label each node with the total weight of the prefix given by the path from the root



```
def prep_init(self,
               verify_key: bytes,
               agg_id: int,
               agg_param: MasticAggParam,
               nonce: bytes,
               correcton_words: list[CorrectionWord],
               input_share: MasticInputShare,
               ) -> tuple[MasticPrepState, MasticPrepShare]:
    (level, prefixes, do_weight_check) = agg_param
    (key, proof_share) = \
        self.expand_input_share(agg_id, input_share)

    # Evaluate the VIDPF.
    (beta_share, out_share, eval_proof) = \
        self.vidpf.eval(agg_id, correction_words, key,
                        level, prefixes, nonce)

    # Query the FLP if applicable.
    verifier_share: Optional[list[F]] = None
    if do_weight_check:
        query_rand = self.query_rand(verify_key, nonce, level)
        verifier_share = \
            self.flp.query(beta_share, proof_share, query_rand)

    prep_share = (eval_proof, verifier_share)
    return (out_share, prep_share)
```